

Compression on MNIST dataset × Welcome to Colab - Colab × +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents

Welcome to Colab!

Getting started

Data science

Machine learning

More resources

Featured examples

Section

```
[2] 3m
import tensorflow as tf
from tensorflow.keras import layers, Model
import numpy as np
import matplotlib.pyplot as plt

# -----
# Load and preprocess MNIST
# -----
(x_train, _), (x_test, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0

x_train = x_train.reshape(-1, 784)
x_test = x_test.reshape(-1, 784)

latent_dim = 2 # size of latent space

# -----
# Encoder
# -----
inputs = layers.Input(shape=(784,))
h = layers.Dense(256, activation='relu')(inputs)
h = layers.Dense(128, activation='relu')(h)
z_mean = layers.Dense(latent_dim)(h)
z_log_var = layers.Dense(latent_dim)(h)

# Sampling layer (reparameterization trick)
def sampling(args):
    z_mean, z_log_var = args
    epsilon = tf.random.normal(shape=(tf.shape(z_mean)[0], latent_dim))
    return z_mean + tf.exp(z_log_var * 0.5) * epsilon
```

Variables Terminal ✓ 00:31 Python 3

Compression on MNIST dataset x Welcome to Colab - Colab x +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents [2] 3m

Welcome to Colab!

Getting started

Data science

Machine learning

More resources

Featured examples

+ Section

```
z = layers.Lambda(sampling)([z_mean, z_log_var])
encoder = Model(inputs, [z_mean, z_log_var, z], name="encoder")

# -----
# Decoder
# -----
latent_inputs = layers.Input(shape=(latent_dim,))
x = layers.Dense(128, activation='relu')(latent_inputs)
x = layers.Dense(256, activation='relu')(x)
outputs = layers.Dense(784, activation='sigmoid')(x)
decoder = Model(latent_inputs, outputs, name="decoder")

# -----
# Custom VAE Model
# -----
class VAE(Model):
    def __init__(self, encoder, decoder, **kwargs):
        super(VAE, self).__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder

    def train_step(self, data):
        if isinstance(data, tuple):
            data = data[0]
        with tf.GradientTape() as tape:
            z_mean, z_log_var, z = self.encoder(data)
            reconstruction = self.decoder(z)

            # Reconstruction loss
            reconstruction_loss = tf.reduce_sum(
                tf.keras.losses.binary_crossentropy(data, reconstruction), axis=-1
            )

```

Variables Terminal ✓ 00:31 Python 3

Compression on MNIST dataset × Welcome to Colab - Colab × +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents [2] 3m

Welcome to Colab!

Getting started

Data science

Machine learning

More resources

Featured examples

+ Section

```
        )
reconstruction_loss = tf.reduce_mean(reconstruction_loss)

# KL divergence loss
kl_loss = -0.5 * tf.reduce_mean(
    1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var)
)

total_loss = reconstruction_loss + kl_loss

grads = tape.gradient(total_loss, self.trainable_weights)
self.optimizer.apply_gradients(zip(grads, self.trainable_weights))

return {
    "loss": total_loss,
    "reconstruction_loss": reconstruction_loss,
    "kl_loss": kl_loss,
}

def test_step(self, data):
    if isinstance(data, tuple):
        data = data[0]
    z_mean, z_log_var, z = self.encoder(data)
    reconstruction = self.decoder(z)

    reconstruction_loss = tf.reduce_sum(
        tf.keras.losses.binary_crossentropy(data, reconstruction), axis=-1
    )
    reconstruction_loss = tf.reduce_mean(reconstruction_loss)

    kl_loss = -0.5 * tf.reduce_mean(
        1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var)
)
```

Variables Terminal ✓ 00:31 Python 3

Compression on MNIST dataset × Welcome to Colab - Colab × +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents [2] 3m

Welcome to Colab!

Getting started

Data science

Machine learning

More resources

Featured examples

Section

```
kl_loss = -0.5 * tf.reduce_mean(  
    1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var)  
)  
  
total_loss = reconstruction_loss + kl_loss  
  
return {  
    "loss": total_loss,  
    "reconstruction_loss": reconstruction_loss,  
    "kl_loss": kl_loss,  
}  
  
# -----  
# Compile and Train  
# -----  
vae = VAE(encoder, decoder)  
vae.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001))  
  
history = vae.fit(  
    x_train,  
    epochs=30,  
    batch_size=256,  
    validation_data=(x_test, None),  
    verbose=1,  
)  
  
# -----  
# Compression (Encoding) and Reconstruction  
# -----  
z_mean, z_log_var, z = encoder.predict(x_test)  
x_decoded = decoder.predict(z)
```

Variables Terminal ✓ 00:31 Python 3

Compression on MNIST dataset × Welcome to Colab - Colab × +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar ... YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents [2] 3m

Welcome to Colab!

Getting started

Data science

Machine learning

More resources

Featured examples

+ Section

```
# -----
# Visualization
# -----
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # Original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap="gray")
    plt.title("original")
    plt.axis("off")

    # Reconstructed
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(x_decoded[i].reshape(28, 28), cmap="gray")
    plt.title("Reconstructed")
    plt.axis("off")
plt.show()

→ Epoch 1/30
235/235 10s 32ms/step - kl_loss: 2.4209 - loss: 69.4352 - reconstruction_loss: 67.0143 - val_kl_loss: 2.0936 - val_loss:
Epoch 2/30
235/235 10s 28ms/step - kl_loss: 1.9075 - loss: 57.1722 - reconstruction_loss: 55.2647 - val_kl_loss: 2.0989 - val_loss:
Epoch 3/30
235/235 6s 27ms/step - kl_loss: 2.0547 - loss: 55.3771 - reconstruction_loss: 53.3224 - val_kl_loss: 2.1453 - val_loss:
Epoch 4/30
235/235 7s 30ms/step - kl_loss: 2.1609 - loss: 54.2927 - reconstruction_loss: 52.1318 - val_kl_loss: 2.2014 - val_loss:
Epoch 5/30
235/235 6s 27ms/step - kl_loss: 2.2781 - loss: 53.3850 - reconstruction_loss: 51.1070 - val_kl_loss: 2.2363 - val_loss:
Epoch 6/30
235/235 11s 32ms/step - kl_loss: 2.3838 - loss: 52.4831 - reconstruction_loss: 50.0993 - val_kl_loss: 2.3522 - val_loss:
Epoch 7/30
```

Variables Terminal ✓ 00:31 Python 3

Compression on MNIST dataset x Welcome to Colab - Colab x +

colab.research.google.com/#scrollTo=Omd8LwiP4czT

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all Copy to Drive RAM Disk

Table of contents

- Welcome to Colab!
- Getting started
- Data science
- Machine learning
- More resources

Featured examples

+ Section

Epoch 23/30
235/235 7s 28ms/step - kl_loss: 2.8707 - loss: 48.1651 - reconstruction_loss: 45.2944 - val_kl_loss: 2.8114 - val_loss:
Epoch 24/30
235/235 8s 35ms/step - kl_loss: 2.8778 - loss: 48.0632 - reconstruction_loss: 45.1854 - val_kl_loss: 2.8724 - val_loss:
Epoch 25/30
235/235 6s 25ms/step - kl_loss: 2.8878 - loss: 47.9787 - reconstruction_loss: 45.0909 - val_kl_loss: 2.8516 - val_loss:
Epoch 26/30
235/235 7s 31ms/step - kl_loss: 2.8995 - loss: 47.9013 - reconstruction_loss: 45.0018 - val_kl_loss: 2.7979 - val_loss:
Epoch 27/30
235/235 6s 26ms/step - kl_loss: 2.9123 - loss: 47.7828 - reconstruction_loss: 44.8706 - val_kl_loss: 2.9044 - val_loss:
Epoch 28/30
235/235 10s 25ms/step - kl_loss: 2.9202 - loss: 47.7358 - reconstruction_loss: 44.8155 - val_kl_loss: 2.8444 - val_loss:
Epoch 29/30
235/235 7s 31ms/step - kl_loss: 2.9242 - loss: 47.6562 - reconstruction_loss: 44.7320 - val_kl_loss: 2.8429 - val_loss:
Epoch 30/30
235/235 9s 26ms/step - kl_loss: 2.9339 - loss: 47.5791 - reconstruction_loss: 44.6452 - val_kl_loss: 2.7903 - val_loss:
313/313 1s 3ms/step
313/313 1s 4ms/step

| Original |
|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
| | | | | | | | | | |
| Reconstructed |
| | | | | | | | | | |

Variables Terminal ✓ 00:31 Python 3

11. Experiments using variational Autoencoder.

Aim:- To implement a variational Autoencoder (VAE) for compressing and reconstructing the MNIST handwritten digit images and to analyze its ability to learn a smooth latent space representation.

objectives:-

1. To understand the working of variational auto encoders as probabilistic extension of standard auto encoders
2. To learn a latent distribution for efficient image compression.
3. To reconstruct MNIST digits from the latent space
4. To visualize and interpret the latent space and reconstructed outputs
5. To compare the performance of VAE with.

Pseudo code:-

1. Import Tensor Flow and MNIST dataset
2. Normalize and flatten the dataset
3. Define encoder:
Input \rightarrow dense(512, relu)
 $\text{dense}(\text{latent_dim}^2) \rightarrow \text{out} \sim \text{mean}(\mu)$
 $\log \text{variance}(\log \sigma^2)$

4. Sampling function:-

$$z = \mu + \exp(0.5 * \log \sigma^2) * E$$

distribution
90%
information.

- 5) Define decoder
 $\text{Input}(z) \rightarrow \text{Dense}(512, \text{relu}) \rightarrow \text{Dense}(784, \text{sigmoid})$
- 6) Define VAE model = Encoder + Sampling + Decoder
- 7) loss = reconstruction loss + KLD divergence loss.
- 8) Train model for 50 epochs
- 9) Visualize original and reconstructed images
- 10) Plot 2D latent space

✓
Sandeep

Result:-

- The variational Autoencoder successfully compressed MNIST digit images into a 2-D latent space.
- The KL divergence ensured a smooth latent distribution enabling interpolation b/w digits.
- Reconstruction loss was $0.02 + 0.03$ depending on training.
- The model demonstrates that VAEs can learn many continuous and structured latent representations useful for data compression and generation.

epoch 1/30 :-

KL loss :- 2.4209 - loss :- 69.4352.
reconstruction - loss :- 67.0143 - Val KL loss :- 2.0936.

epoch 10/30 :-

KL loss :- 2.5976 - loss :- 50.5912.
reconstruction - loss :- 47.9936 - Val KL loss :- 2.6866

epoch 20/30 :-

KL loss :- 2.8249 - loss :- 48.5557 -

reconstruction - loss :- 45.9307 - Val KL loss :- 2.7994

epoch 30/30 :-

KL loss :- 2.9339 - loss :- 47.5791 -

reconstruction - loss :- 44.6643 - Val loss :- 2.2903

Original

7 2 0 0 0 4 0

reconstructed

7 2 0 0 0 4 0