

Ex:- 14:- Implement a pre-trained CNN model as Feature Extractor using Transfer learning models.

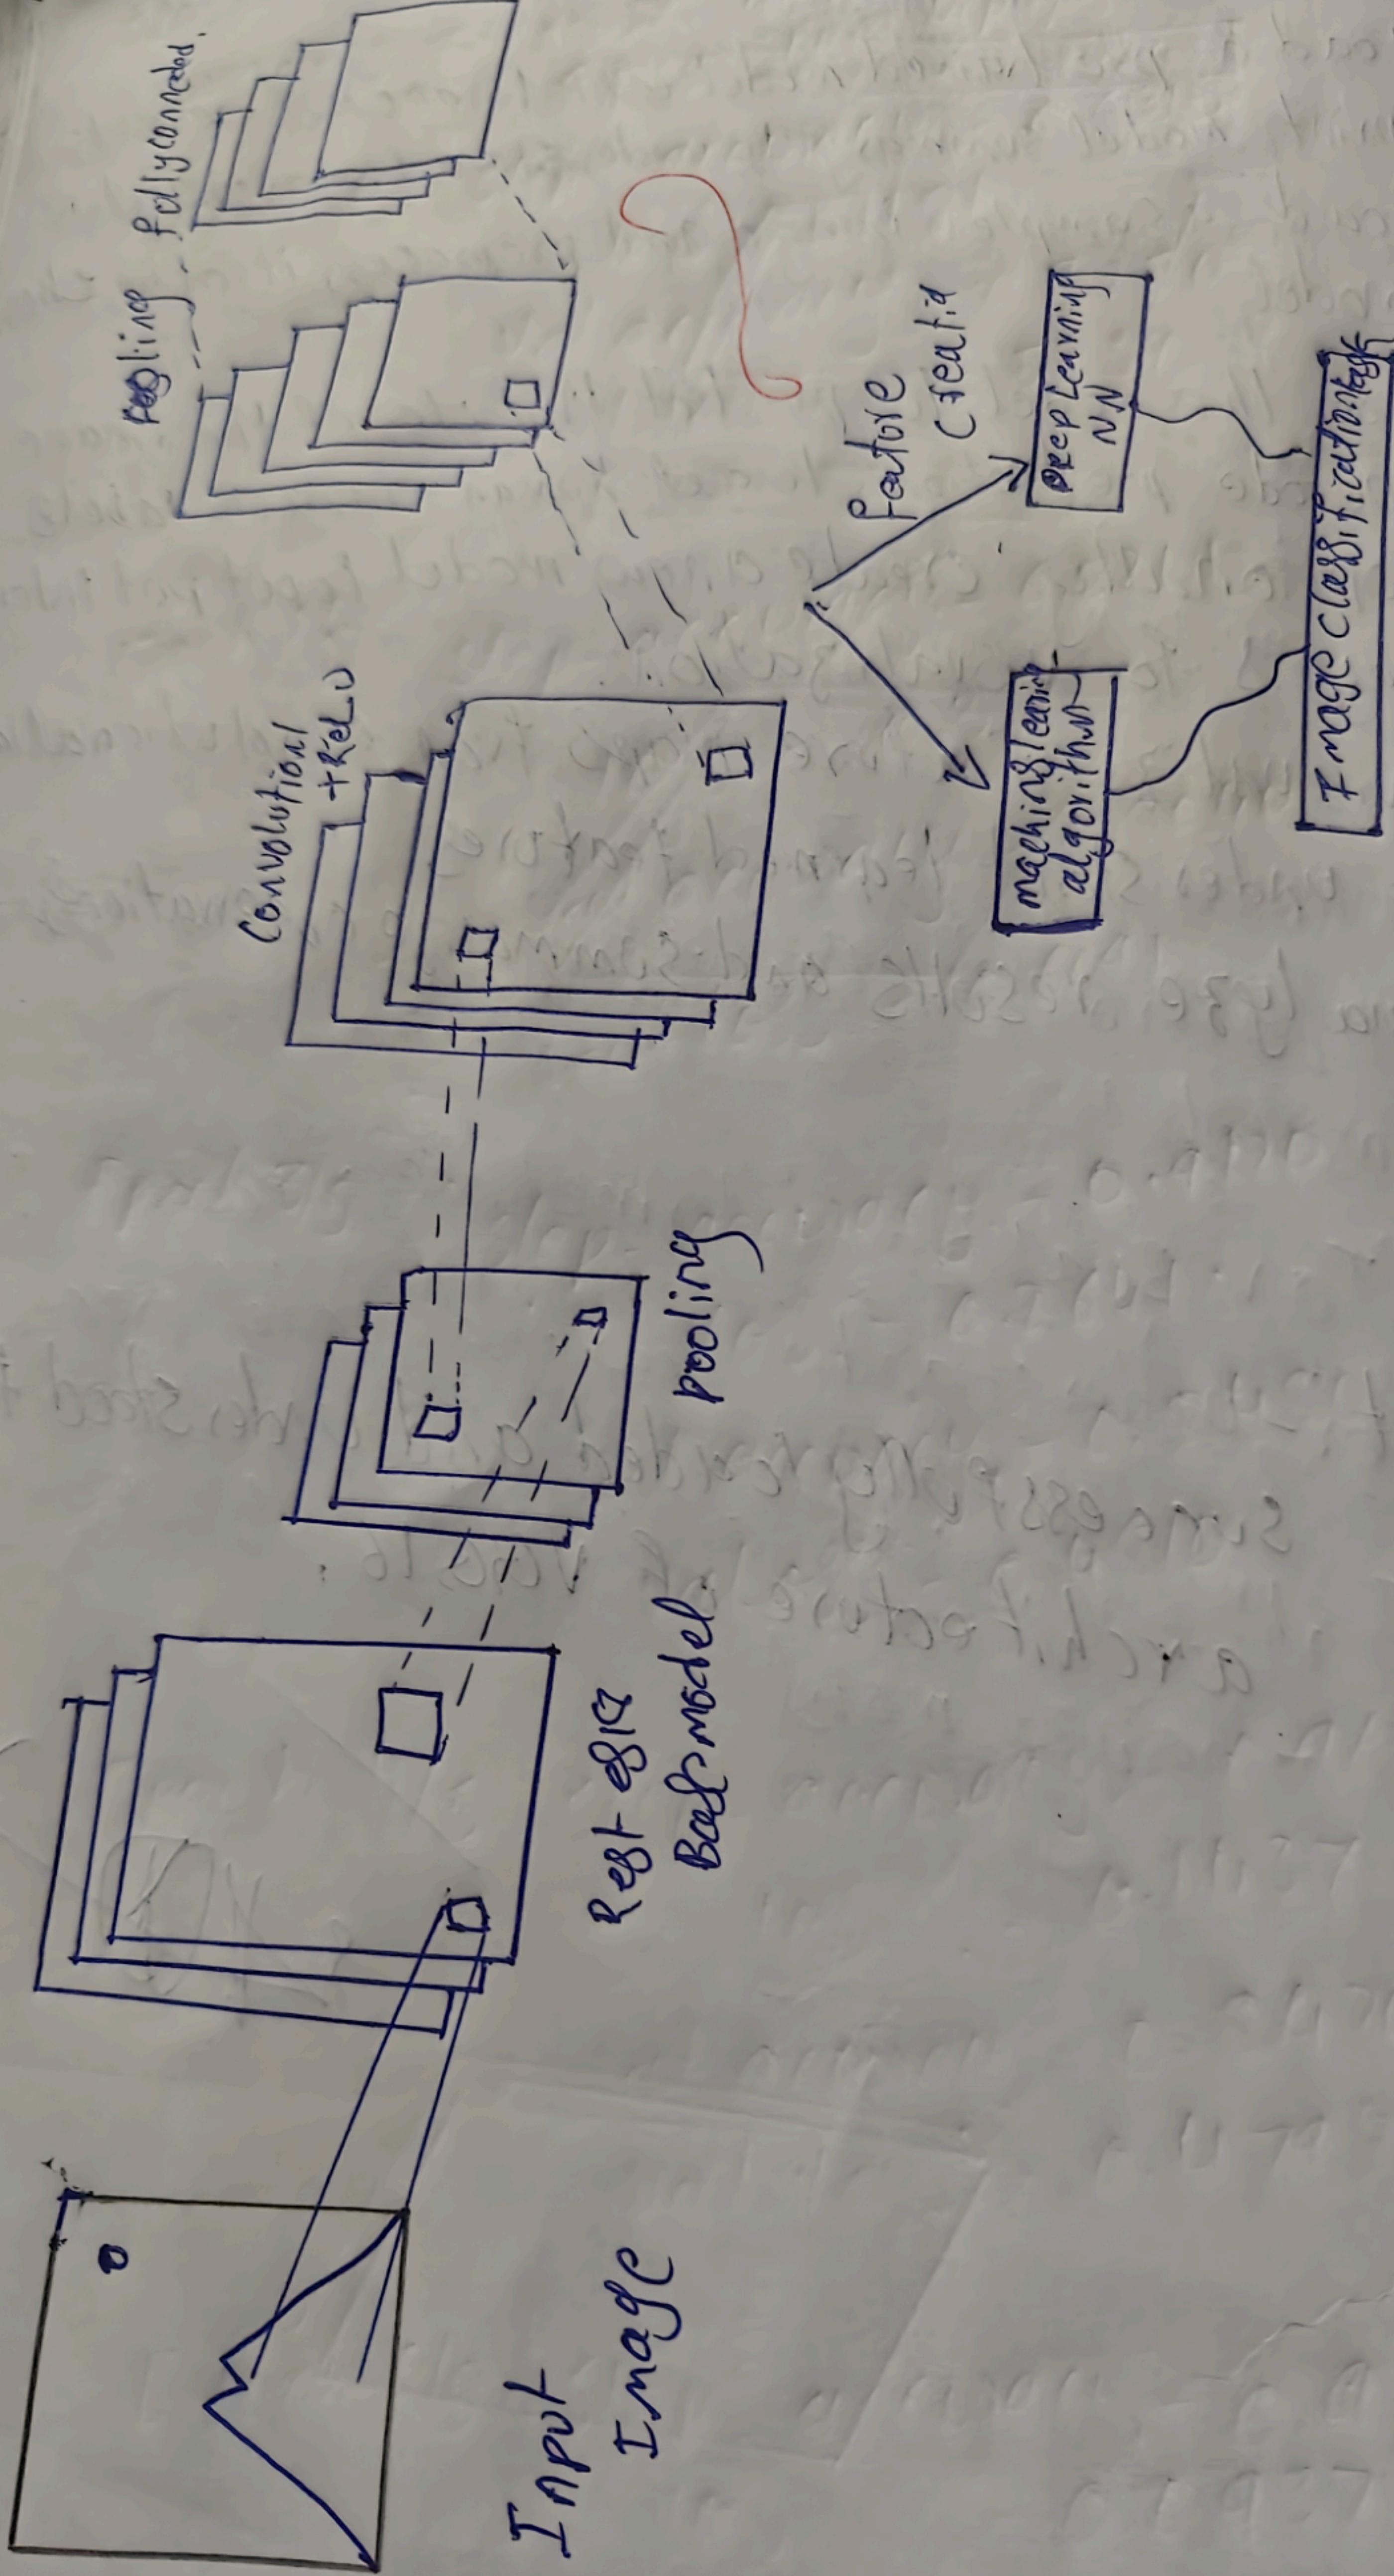
Aim:- To implement a memory-efficient pre-trained CNN as a feature extractor for image classification on CIFAR-10.

Objectives:-

1. Load a pre-trained MobileNet V2 model trained on ImageNet.
2. Freeze all pre-trained layers to prevent weight updates.
3. Resize and preprocess CIFAR-10 images to match MobileNet V2 input requirements.
4. Build a classifier on top of mobile net V2 using Global Average Pooling 2D and Dense layers.
5. Train the classifier directly on batches of images avoiding precomputing and storing features.

(6. Evaluate the model (a))

Transfer learning.



Pseudo code:-

1. Import tensorflow, keras, numpy.
2. Load CIFAR-10 dataset
3. Resize images to mobile net V2 input size (96x96).
4. preprocess images using preprocess input
5. Load pre trained layers
6. Freeze pre trained layers
7. Build model
 - a. mobile net V2 base
 - b. Global Average pooling 2D.
 - c. Dense (10, activation = softmax)
 - d. Dense (28, activation = relu)
8. compile model (optimizer = adam, loss = sparse_categorical_crossentropy)
9. Train model directly on batches of images.
10. Evaluate model on test data.

Result:- successfully implemented mobile Net V2 as Feature Extractor using transfer learning model

✓
YES

Out put:-
1. memory efficient :- No precomputed feature arrays.
2. model runs on low - Ram laptops.

Epoch 1/5 Step - accuracy :- 0.6315
loss :- 1.4181

Val accuracy :- 0.8250
val loss :- 0.5289

Epoch 2/5 Step - accuracy :- 0.8761

loss :- 0.3891

Val accuracy :- 0.8040

Val loss :- 0.6041

Epoch 3/5 Step - accuracy :- 0.9120

loss :- 0.2641

Val accuracy :- 0.8480

Val loss :- 0.4810

Epoch 4/5 Step - accuracy :- 0.9610

loss :- 0.1437

Val accuracy :- 0.8470

Val loss :- 0.4798

Epoch 5/5 Step - accuracy :- 0.9757

loss :- 0.0937

Val accuracy :- 0.8320

Val loss :- 0.5536.

Pseudocode
1. Import
2. Load
3. Reshape
4. Preprocess
5. Load
6. Fine-tune
7. Build

8.

9.

R

Welcome to Colab - Colab x Compression on MNIST dataset x +

colab.research.google.com/#scrollTo=1a5MOUjTid8e

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all Copy to Drive RAM Disk

Table of contents + Section

```
[3] ✓ 36s
import tensorflow as tf
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Use small subset for low RAM
x_train, y_train = x_train[:5000], y_train[:5000]
x_test, y_test = x_test[:1000], y_test[:1000]

# Resize images to 96x96 and preprocess
x_train_resized = tf.image.resize(x_train, (96,96))
x_test_resized = tf.image.resize(x_test, (96,96))
x_train_preprocessed = preprocess_input(x_train_resized)
x_test_preprocessed = preprocess_input(x_test_resized)

# Load MobileNetV2 and freeze layers
base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(96,96,3))
base_model.trainable = False

# Build memory-efficient model
model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')]
```

Variables Terminal ✓ 03:31 T4 (Python 3)

Welcome to Colab - Colab Compression on MNIST dataset

colab.research.google.com/#scrollTo=1a5MOUjTid8e

Gmail Maps News Translate Disney+ Hotstar YouTube All Bookmarks

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive RAM Disk

Table of contents + Section

```
[3] [3]
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train model directly
history = model.fit(x_train_preprocessed, y_train, epochs=5, batch_size=32,
                     validation_data=(x_test_preprocessed, y_test))

# Evaluate
loss, accuracy = model.evaluate(x_test_preprocessed, y_test)
print("Test Accuracy:", accuracy)
```

Epoch 1/5
157/157 22s 76ms/step - accuracy: 0.6315 - loss: 1.1181 - val_accuracy: 0.8250 - val_loss: 0.5289
Epoch 2/5
157/157 2s 12ms/step - accuracy: 0.8711 - loss: 0.3891 - val_accuracy: 0.8040 - val_loss: 0.6091
Epoch 3/5
157/157 2s 11ms/step - accuracy: 0.9120 - loss: 0.2641 - val_accuracy: 0.8480 - val_loss: 0.4810
Epoch 4/5
157/157 2s 11ms/step - accuracy: 0.9610 - loss: 0.1437 - val_accuracy: 0.8470 - val_loss: 0.4798
Epoch 5/5
157/157 2s 15ms/step - accuracy: 0.9757 - loss: 0.0937 - val_accuracy: 0.8320 - val_loss: 0.5536
32/32 0s 10ms/step - accuracy: 0.8478 - loss: 0.5003
Test Accuracy: 0.8320000171661377

Variables Terminal ✓ 03:31 T4 (Python 3)